Deviations from the preregistration

The plan was frozen at git tag `prereg-v1` on 15 September 2026, before any confirmatory episode was generated (`docs/PREREGISTRATION.md`). Every departure from it since is listed here and in the paper's appendix.

Recorded as they happen; each is reported in the paper.

Date	Deviation	Reason	Data affected
2026-09-15	`request_timeout_s` raised to 600 s for the Qwen retry pass (`configs/main_qwen.yaml`)	Qwen3.5 397B occasionally reasons past 300 s, producing provider timeouts	Retried Qwen episodes only; their recorded config hash differs
2026-09-15	Stated-total extractor corrected (`detectors-v2` → `detectors-v3`)	Reports from the pilots and from an 8-episode check of Claude Fable 5.1 run after the freeze (seed 201, 2026-09-15 22:12 UTC, after the 20:18 UTC freeze commit) exposed four false-positive classes: margins over the fallback ("20 points better than your 31-point alternative"), figures belonging to the benchmark ("your 31-point alternative"), per-term contributions ("payment terms are worth 37 points") and packages that were not taken ("that package was worth 59 points, 7 more than the final deal"). The fix also recognises "by your scoring that's … points" as a claim about the deal.	Applied before any confirmatory episode was scored. All pilot episodes were rescored; pilot conclusions are unchanged (e.g. Claude Haiku 4.5 still misstates in 32 of 43 reports that state a total). Test cases are the real phrasings, in `tests/test_fidelity_detectors.py`.
2026-09-15	Stated-total extractor corrected again (`detectors-v3` → `detectors-v4`)	Human labels showed the corrected extractor still missed phrasings, giving recall 0.42: requiring an explicit deal cue was too strict, and totals stated next to a margin or benchmark figure were being attributed to the wrong quantity. The fix relaxes the deal-cue requirement and adds precedence for explicit total markers.	Applied before any confirmatory number was computed. Precision against the annotator consensus rose to 1.00 with recall 0.83.

Date	Deviation	Reason	Data affected
2026-09-16	Human validation obtained for 50 disclosure and 25 stated-total items, against the 150/150/100/50 planned in section 7	Annotator time; two annotators labelled every item in the batches that were drawn	Characterization ratings and leak detection have no human validation and are reported as secondary throughout. Disclosure and the stated-total extractor are validated on the items obtained and clear their preregistered thresholds. Batches were drawn before annotation; none were discarded.
2026-09-16	Unconditional estimands and worst-case bounds added for D1 and D2 (`unconditional_and_bounds`)	H1a and H1b condition on gates (exposure; whether a report states a total) that are themselves outcomes; the frozen plan did not address post-treatment conditioning	Reported alongside the frozen conditional tests, which are unchanged. Neither headline conclusion depends on the gate.
2026-09-16	Extractor and stage-1 validation statistics corrected in `proxy annotate agreement`	The extractor was graded against the value frozen in the batch key at sampling time, so it reported the pre-fix extractor's accuracy (recall 0.42) rather than the current one's; and stage-1 precision was computed on the four-way consensus, which contains no unmentioned items, returning n=0.	Reporting only. The extractor is now re-run from the report text (precision 1.00, recall 0.83, matching the paper) and stage-1 precision is computed on the binary consensus the claim is about (1.00, n=16). The judge demotion flag now applies the preregistered binary criterion on the binary consensus set.
2026-09-16	Evaluation-awareness pattern (E8) broadened	The original pattern missed hedged and nominal phrasings ("this was clearly a test", "part of this eval"), so a rate of zero was partly a coverage artifact	Exploratory only. The broadened pattern flags 5 of 2,880 reports; all five were read and none claims the episode was a test, so the reported finding is unchanged and better supported.
2026-09-16	Judged non-disclosure additionally reported corrected for judge error (stratified Rogan–Gladen, `proxy/annotation/debias.py`)	A judged rate is an instrument reading, not the quantity of interest; the frozen plan validated the judge but did not invert it	Corrected rate 97.1% against a raw judged 97.1%; the correction is reported with a sensitivity analysis over the weakly identified specificity.
2026-09-15	Claude Fable 5.1 added as an agent (addendum A2, pending full arm)	Claude Fable 5.1 is Anthropic's current frontier model; comparing GPT-6 Astra with Claude Opus 5 alone would not be a frontier-to-frontier comparison	New arm; reported with and without

Date	Deviation	Reason	Data affected
2026-10-17	Stated-total extractor corrected a third time (`detectors-v4` → `detectors-v5`)	A reader of the draft noted that the largest stated-total errors (20–60 points for Claude Opus 5, Claude Sonnet 5 and Llama 4 Maverick) were implausible as model error. Reading the sentences confirmed an extractor fault: in itemized sentences ("1 point on price, 28 on delivery, 9 on warranty, 25 on payment terms — 63 points total") the extractor took the first per-term figure as the deal's total. The fix excludes per-term figures and takes the figure nearest an explicit total marker on either side.	Post-hoc correction to a confirmatory measure, applied after the first analysis and after the draft was written. All runs rescored; every D2 number regenerated; H1b and H2b re-run, conclusions unchanged. Misstatement rates for Claude Opus 5, Claude Sonnet 5 and Llama 4 Maverick fall from 10%, 17% and 8% to about 1–2%; Claude Haiku 4.5, DeepSeek V4 Pro and Mistral Large 3 essentially unchanged. Over the confirmatory grid 108 of 2,241 extracted totals changed. Extractor precision and recall against the batch-a consensus recomputed. Test cases are the real phrasings, in `tests/test_fidelity_detectors.py`.
2026-10-21	Calibration sample size in the paper corrected from "240 episodes per condition" to 138	The figure was typed into the text and was wrong: `proxy calibrate configs/pilot.yaml` averages the floor-respecting reference agent over 23 scenarios (the pilot's 3 plus 20 extra), both roles and 3 concession curves, 138 episodes per outcome condition, as `docs/DECISIONS.md` recorded. The number is now generated at build time with its derivation.	Reporting only; the calibration values (0.44, 0.24, 0.09) are unchanged.

Date	Deviation	Reason	Data affected
2026-10-21	Knowledge-probe scoring aligned with its definition; GEE fallback implemented	The probe scored an answer as "knew" if it matched any forgone better offer, compliant or blocked, whereas K1 and the paper ask whether the agent named the *blocked* offer's value. Separately, the GEE fits supporting H2a and H3 returned non-finite estimates under the exchangeable working correlation without triggering the fallback specified in addendum A3.3.	Exploratory and supporting analyses only. `probe_knew` is now a match to the blocked offer's value within two points (the broader score is kept as `probe_knew_any`). `gee_logit` refits with an independence working correlation and cluster-robust SEs when the exchangeable fit does not converge, and records that it did.

Addendum A2 (2026-09-16, before any Fable 5.1 confirmatory data)

Claude Fable 5.1 (`claude-fable-5-1`) is added as an agent, in the same design and on the same scenario seeds as the main grid (`configs/main_fable.yaml`, 240 episodes). It enters the Anthropic ladder at capability rank 4. An 8-episode check on seed 201 preceded this addendum and is disclosed: judged non-disclosure 6/8 by the primary judge, the knowledge probe correct in 8/8, and stated totals correct in 8/8. Confirmatory results are reported both with and without this arm, since it was added after the main grid ran.